

Integrating Elasticsearch into Analytics Workflows

REV2
May 23, 2019

STEPHANIE KIRMER



@data_stephanie

github.com/skirmer/elastic_analytics

AGENDA

**Introducing Elasticsearch
Libraries for R and Python
Querying and Filtering
Summarizing
Further Reading**

^

Introducing Elasticsearch



Overview

- Part of a family of data storage options called **NoSQL**
 - **Not the same as tabular or SQL style data storage**
- Optimized for fast and **powerful searching**
- **Scales to "big data"**— but usable for small projects
- **Open source tool**

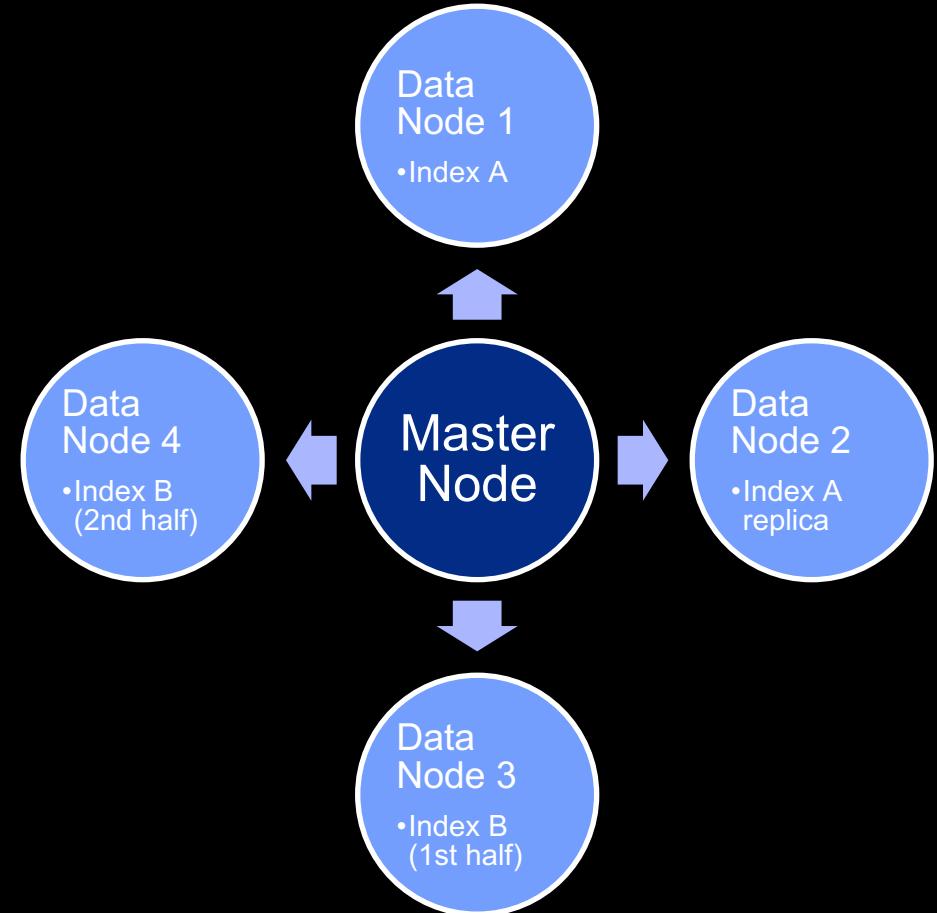
*Sometimes abbreviated "ES"



Visualizing Elasticsearch Storage

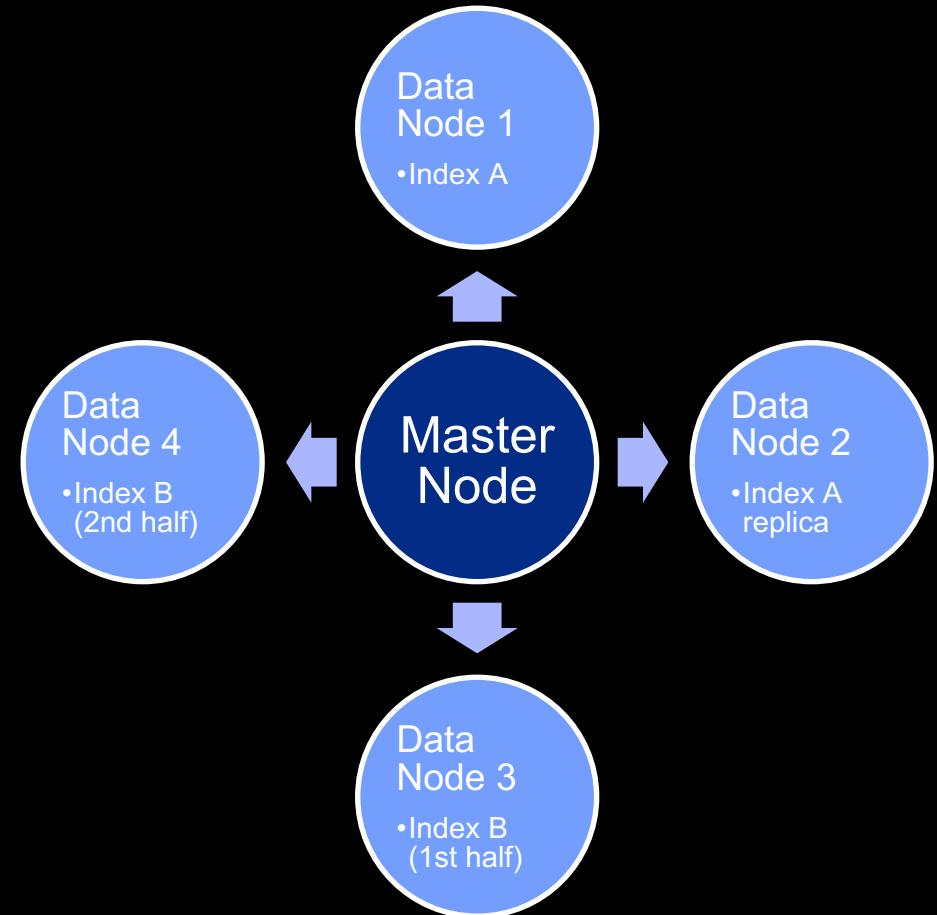
Data Architecture

- Data is divided into **indices**
- Distributed
- **Indices :**
 - Are user-defined groupings of data with some commonality
 - can live on one node, or
 - can be "sharded" and broken across nodes
 - can be duplicated on different nodes



Visualizing Elasticsearch Storage

- **Cluster** = group of nodes
- **Master Node** = central manager
 - Manages indices
 - Tracks and organizes nodes
 - Decides which shards to allocate to nodes
- **Data Node** = where data is kept/handled
 - Search, aggregation functions



A Sidebar About “Search”

Searching	Not Searching (Querying, for example)
Accommodates uncertainty, ambiguity	Requires precise, specific, clear requests
Tries to help you figure out what you need	Extremely literal
Example: Googling “data storage ideas”	Example: Typing https://www.elastic.co/ in browser

Elasticsearch is designed to be great for searching

- Built on top of a technology called Lucene from Apache – Java only
- Allows easy API access to Lucene without Java

Elasticsearch (thanks to Lucene) can handle:

- Typos
- Wildcards
- Fuzzy matching
- Updating and searching simultaneously

Tabular Data vs Document-Based Data

NoSQL is a different paradigm for thinking about data.

Data courtesy of Annie Millerbernd of the San Antonio Express-News. You can learn more about it and see the original dataset here: <https://data.world/amillerbernd/ut-system-post-grad-earnings>

	institution_id	institution_name	deglevl_code	deglevel	degcip_4dig	ciptitle
1	3599	UNIVERSITY OF TEXAS - RIO GRANDE VALLEY	3	Baccalaureate	301	NATURAL RESOURCES CONSERVATION AND RESEARCH
2	3599	UNIVERSITY OF TEXAS - RIO GRANDE VALLEY	3	Baccalaureate	301	NATURAL RESOURCES CONSERVATION AND RESEARCH
3	3599	UNIVERSITY OF TEXAS - RIO GRANDE VALLEY	3	Baccalaureate	501	AREA STUDIES
4	3599	UNIVERSITY OF TEXAS - RIO GRANDE VALLEY	3	Baccalaureate	501	AREA STUDIES
5	3599	UNIVERSITY OF TEXAS - RIO GRANDE VALLEY	3	Baccalaureate	501	AREA STUDIES
6	3599	UNIVERSITY OF TEXAS - RIO GRANDE VALLEY	3	Baccalaureate	501	AREA STUDIES
7	3599	UNIVERSITY OF TEXAS - RIO GRANDE VALLEY	3	Baccalaureate	501	AREA STUDIES
8	3599	UNIVERSITY OF TEXAS - RIO GRANDE VALLEY	3	Baccalaureate	501	AREA STUDIES

```
{"_index":"utexas",
"_type":"data",
"_id":"AWbU6WJiWX1fgzrfh4p1",
"_score":1.0,
"_source":
  {"institution_id":3599,
  "institution_name":"UNIVERSITY OF TEXAS - RIO GRANDE VALLEY",
  "deglevl_code":3,
  "deglevel":"Baccalaureate",
  "degcip_4dig":901,
  "ciptitle":"COMMUNICATION AND MEDIA STUDIES",
  "grad_cohort":2007,
  "grad_cohort_label":"2007-2009",
  "year_postgrad":1,
  "p25_earnings":26518.57,
  "p50_earnings":42166.31,
  "p75_earnings":50439,
  "system":"utsys",
  "cellcount":70}}
```

Why Use Elasticsearch?



Safe

- Copying your data easily and conveniently (via replicas) = if a node fails, your data is safe

Fast

- ES can search in parallel on multiple nodes and replicas, and find your data faster

Scalable

- Once you establish your ES database, you can add nodes and allow your database to grow

Flexible

- Robust search helps you by discerning typos, ranking results, parsing text, and more

Open Source

- Free to use at small scale, substantial documentation, community support

Among other reasons!

Query Language Crash Course

Elastic Query DSL (domain specific language): a JSON-style syntax built to interact with ES databases.

Why use query language?

Consistency across interfaces and media

Precision and power in search, filtering, and aggregating – ES was built to work with this.

Downsides?

It's sometimes hard to work with – idiosyncratic rules of syntax.



Handy docker tips

docker ps to check your containers

docker kill [name] to hard shutdown your containers

11

Follow Along!

When you see this arrow,
you can try it out yourself!

System Requirements:

- Docker installed and running
- Repository ready: `git clone https://github.com/skirmmer/elastic_analytics.git`

Setup Steps (see the README for commands to copy/paste)

- Get into the top level of the cloned repo
- At Terminal:
 1. `./supporting_materials/setup_texas.sh 5.5`
 2. `curl -X POST 'http://localhost:9200/utexas/_bulk' -H 'Content-Type: application/json' --data-binary @supporting_materials/ut_data.json`

Start up R/Rstudio or Python as you prefer, and run further commands from there.



Libraries for R and Python

Choosing the right tool for your
needs

Library Characteristics

Library	Returns	Query Language	Supports Authentication	R	Python
uptasticsearch	Tabular	Required	✗	●	●
elastic	JSON	Supported, not required	●	●	✗
elasticsearch-py	JSON	Supported, not required	●	✗	●

For python: pip install [library name] (use python3)

For R: install.packages("library_name")

KEY CONSIDERATIONS

Secure Authentication

Do you need to securely log in?

Output Format

Do you mind handling JSON output?

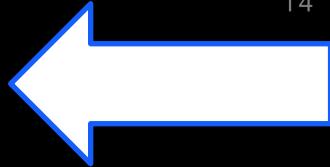
Query Construction

Is writing query language a barrier?



R Options

14



Uptasticsearch (R)

```
test_up <- uptasticsearch::es_search(  
  es_host = "http://localhost:9200"  
  , es_index = "utexas"  
  , query_body = query_string  
  , size = 10)
```

```
query_string <- '{"query": {"match_all":{}}}'
```

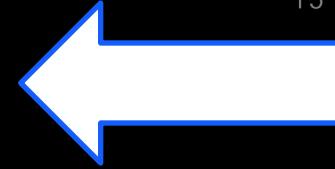
Elastic (R)

```
elastic::connect(es_host =  
  "http://localhost:9200")  
  
test_e <- elastic::Search(index =  
  "utexas"  
  , body = query_string  
  , size = 10  
  , raw = TRUE)  
  
test_e2 <-  
  jsonlite::fromJSON(test_e)$hits$hits
```

Non-Query Language Option:

```
test_e <- elastic::Search(  
  index = "utexas"  
  , q = "grad_cohort:*"  
  , size = 10  
  , raw = TRUE)
```

Python Options



Uptasticsearch (Py)

```
import json
import uptasticsearch

Uptasticsearch.es_search(
    es_host="http://localhost:9200",
    query_body=query_string,
    es_index="utexas"
)
```

```
query_dict = {"query": {"match_all": {}}}
query_string = '{"query": {"match_all": {}}}'
```

Elasticsearch-py (Py)

```
from elasticsearch import Elasticsearch

es = Elasticsearch(['http://localhost:9200'])

res = es.search(
    index="utexas",
    body= query_dict
)

res['hits']['hits']
```

Tip! quickly format using
from pandas.io.json
import json_normalize

Non-Query Language Option (Elasticsearch_dsl):

```
from elasticsearch_dsl import Search

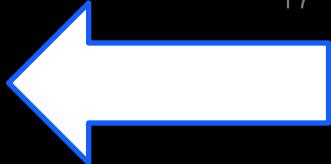
res2 = Search(using = es).query("match",
    _index = 'utexas').execute()

res2.to_dict()['hits']['hits']
```



Querying and Filtering

Get what you need out of your
database



Identifying Available Fields

R:

```
uptasticsearch::get_fields(es_host = "http://localhost:9200",  
es_indices = "utexas")
```

At Command Line:

```
curl http://localhost:9200/utexas/_mapping > fields.json
```

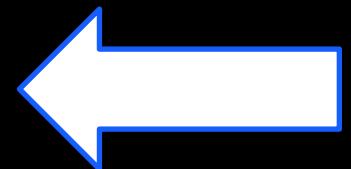
Constructing a Basic Query

Return all records :

```
{  
  "query": { "match_all": {} }  
}
```

```
query_dict = {"query": {"match_all": {}}}  
query_string = '{"query": {"match_all": {}}}'
```

Remember this from earlier!
All the queries we look at can be
passed to R or Python this way.



Constructing a Basic Query

Return all records :

```
{  
  "query" : { "match_all": {} }  
}
```

Constructing a Basic Query

Match one field :

```
{  
  "query": { "match": { "ciptitle.raw": "COMPUTER SCIENCE" } }  
}
```

Tip: Yellow text means something new has been added to the query

Result Sample

	{'query': {'match': {'ciptitle.raw': 'COMPUTER SCIENCE'}}}									
	_id	_index	_score	_source.cellcount	_source.ciptitle	_source.degcip_4dig	...	_source.p25_earnings	_source.p50_earnings	_source.p75_earnings
0	AWq8Wt1a9WDJ8JPhaXL7	utexas	6.612041	-1	COMPUTER SCIENCE	1107	...	NaN	NaN	NaN
1	AWq8Wt0n9WDJ8JPhaXCx	utexas	5.731027	47	COMPUTER SCIENCE	1107	...	43797.75	56302.13	65343.96
2	AWq8Wt0n9WDJ8JPhaXCy	utexas	5.731027	53	COMPUTER SCIENCE	1107	...	47552.69	61813.50	79846.65
3	AWq8Wt1a9WDJ8JPhaXL-	utexas	5.731027	-1	COMPUTER SCIENCE	1107	...	NaN	NaN	NaN
4	AWq8Wt0n9WDJ8JPhaXCr	utexas	5.578471	-1	COMPUTER SCIENCE	1107	...	NaN	NaN	NaN
5	AWq8Wt1a9WDJ8JPhaXL9	utexas	5.578471	-1	COMPUTER SCIENCE	1107	...	NaN	NaN	NaN
6	AWq8Wt1a9WDJ8JPhaXL_-	utexas	5.578471	-1	COMPUTER SCIENCE	1107	...	NaN	NaN	NaN
7	AWq8Wt1a9WDJ8JPhaXMC	utexas	5.578471	44	COMPUTER SCIENCE	1107	...	43425.67	53358.00	63429.86
8	AWq8Wt0n9WDJ8JPhaXCs	utexas	5.192957	-1	COMPUTER SCIENCE	1107	...	NaN	NaN	NaN
9	AWq8Wt0n9WDJ8JPhaX Ct	utexas	5.192957	-1	COMPUTER SCIENCE	1107	...	NaN	NaN	NaN

Things to notice:

- Only Computer science is shown
- NaNs are present in earnings, cellcount sometimes is -1
(this means redacted due to small group size)

Constructing a Basic Query

Match one field AND Greater Than one field :

```
{  
  "query":  
    { "bool" : {  
      "must" : [ { "match": { "ciptitle.raw": "COMPUTER SCIENCE" } }  
                , { "range" : { "cellcount" : { "gt" : 0 } } } ]  
    } }  
}
```

Result Sample

	_id	_index	_score	_source.cellcount	_source.ciptitle	_source.degcip_4dig	...	_source.p25_earnings	_source.p50_earnings
0	AWq8Wt0n9WDJ8JPhaXcx	utexas	6.731027	47	COMPUTER SCIENCE	1107	...	43797.75	56302.13
1	AWq8Wt0n9WDJ8JPhaXcy	utexas	6.731027	53	COMPUTER SCIENCE	1107	...	47552.69	61813.50
2	AWq8Wt1a9WDJ8JPhaXmc	utexas	6.578471	44	COMPUTER SCIENCE	1107	...	43425.67	53358.00
3	AWq8Wt0n9WDJ8JPhaXcu	utexas	6.192957	53	COMPUTER SCIENCE	1107	...	33636.58	41193.17
4	AWq8Wt0n9WDJ8JPhaXcw	utexas	6.192957	40	COMPUTER SCIENCE	1107	...	46477.67	57318.00
5	AWq8Wt1a9WDJ8JPhaXmb	utexas	6.192957	47	COMPUTER SCIENCE	1107	...	44541.92	57588.90
6	AWq8Wt0n9WDJ8JPhaXcv	utexas	6.174845	61	COMPUTER SCIENCE	1107	...	44244.25	58401.60
7	AWq8Wt0n9WDJ8JPhaXcz	utexas	6.174845	55	COMPUTER SCIENCE	1107	...	28598.25	44169.83
8	AWq8Wt1a9WDJ8JPhaXma	utexas	6.174845	36	COMPUTER SCIENCE	1107	...	73903.50	85285.43

Things to notice:

- Only Computer science is shown
- Cellcount values are all above 0 - no more NaNs or -1

Constructing a Basic Query

Match two fields AND Greater Than one field :

```
{  
  "query":  
    { "bool" : {  
      "must" : [ { "match": { "ciptitle.raw": "COMPUTER SCIENCE" } }  
                , { "match": { "institution_id": "3599" } }  
                , { "range" : { "cellcount" : { "gt" : 0 } } } ]  
    } }  
}
```

Result Sample

	<code>_id</code>	<code>_index</code>	<code>_score</code>	<code>_source.cellcount</code>	<code>_source.ciptitle</code>	<code>_source.degcip_4dig</code>	...	<code>_source.p25_earnings</code>
0	Awq8Wt0n9WDJ8JPhaXCx	utexas	7.731027	47	COMPUTER SCIENCE	1107	...	43797.75
1	Awq8Wt0n9WDJ8JPhaXCy	utexas	7.731027	53	COMPUTER SCIENCE	1107	...	47552.69
2	Awq8Wt0n9WDJ8JPhaXCu	utexas	7.192957	53	COMPUTER SCIENCE	1107	...	33636.58
3	Awq8Wt0n9WDJ8JPhaXCw	utexas	7.192957	40	COMPUTER SCIENCE	1107	...	46477.67
4	Awq8Wt0n9WDJ8JPhaXCv	utexas	7.174845	61	COMPUTER SCIENCE	1107	...	44244.25
5	Awq8Wt0n9WDJ8JPhaXCz	utexas	7.174845	55	COMPUTER SCIENCE	1107	...	28598.25

<code>_source.p50_earnings</code>	<code>_source.p75_earnings</code>	<code>_source.system</code>	<code>_source.year_postgrad</code>	<code>_type</code>
56302.13	65343.96	utsys		1 data
61813.50	79846.65	utsys		5 data
41193.17	51356.40	utsys		1 data
57318.00	89080.00	utsys	10	data
58401.60	72430.97	utsys	5	data
44169.83	60239.68	utsys		1 data

Things to notice:

- Only Computer science is shown
- Cellcount values are all above 0
- Only UT-Rio Grande Valley is shown

match_phrase

Match a set of words all together.

exists

Supply a field, returns documents that have at least one non-null value in the original field.

wildcard

Pass a string with a wildcard anywhere – but be careful, it can be a slow search!

filter

Just like "must" except without scoring – we'll talk about this in a moment.

must_not

Instead of "must" – use to omit records with a word or phrase.

This is just a small sample- ES query language offers many very powerful search options!

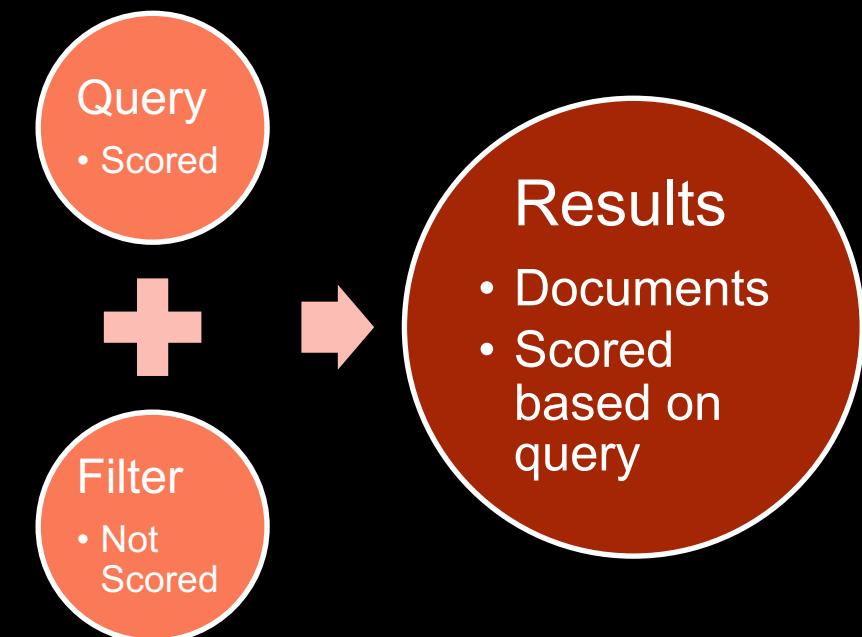
Query vs Filter

Scoring Results

ES queries can provide a **numeric score** indicating how well the document meets the criteria given.

When you use "**query**" at the beginning of the query, you get a score returned alongside your results.

When you use "**filter**", Elasticsearch does not score the results on the given criteria.



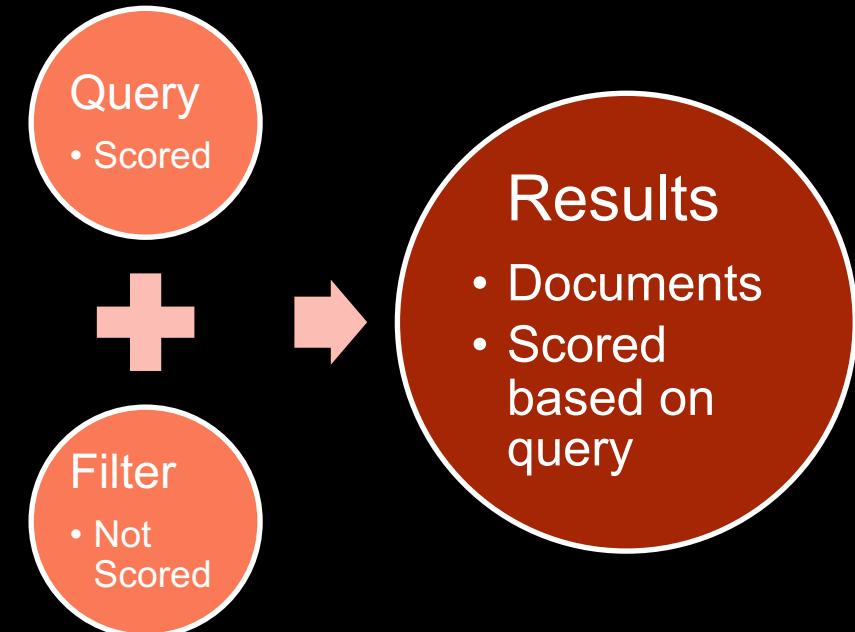
Query vs Filter

Example

```
{
  "query": {
    "bool": {
      "must": [
        { "match": { "ciptitle.raw": "AREA STUDIES" } },
        { "match": { "deglevel": "Baccalaureate" } }
      ],
      "filter": [
        { "match": { "institution_id": "3599" } }
      ]
    }
  }
}
```

Do we want scores returned for this search? Yes.

Do we want the scores to include this criterion? NO.



Query vs Filter

Example

Same query, first with two criteria scored (1 in filter) and second with only 1 criterion scored.

	<code>_id</code>	<code>_index</code>	<code>_score</code>	<code>_source.cellcount</code>	<code>_source.ciptitle</code>	<code>_source.degcip_4dig</code>	...	"match": {
0	AWq8Wt0n9WDJ8JPhaXCJ	utexas	5.474233	-1	AREA STUDIES	501	...	"deglevel":
1	AWq8Wt0n9WDJ8JPhaXCB	utexas	5.168861	-1	AREA STUDIES	501	...	"Baccalaureate"
2	AWq8Wt0n9WDJ8JPhaXCI	utexas	5.168861	-1	AREA STUDIES	501	...	In query (scored)
3	AWq8Wt0n9WDJ8JPhaXCE	utexas	5.154995	-1	AREA STUDIES	501	...	
4	AWq8Wt0n9WDJ8JPhaXCG	utexas	5.154995	-1	AREA STUDIES	501	...	
5	AWq8Wt0n9WDJ8JPhaXCC	utexas	4.792192	-1	AREA STUDIES	501	...	
6	AWq8Wt0n9WDJ8JPhaXCH	utexas	4.792192	-1	AREA STUDIES	501	...	
7	AWq8Wt0n9WDJ8JPhaXCD	utexas	4.733091	-1	AREA STUDIES	501	...	
8	AWq8Wt0n9WDJ8JPhaXCF	utexas	4.733091	-1	AREA STUDIES	501	...	
[9 rows x 18 columns]								
	<code>_id</code>	<code>_index</code>	<code>_score</code>	<code>_source.cellcount</code>	<code>_source.ciptitle</code>	<code>_source.degcip_4dig</code>	...	"match": {
0	AWq8Wt0n9WDJ8JPhaXCJ	utexas	5.145704	-1	AREA STUDIES	501	...	"deglevel":
1	AWq8Wt0n9WDJ8JPhaXCB	utexas	4.843724	-1	AREA STUDIES	501	...	"Baccalaureate"
2	AWq8Wt0n9WDJ8JPhaXCI	utexas	4.843724	-1	AREA STUDIES	501	...	In filter (not scored)
3	AWq8Wt0n9WDJ8JPhaXCE	utexas	4.813467	-1	AREA STUDIES	501	...	
4	AWq8Wt0n9WDJ8JPhaXCG	utexas	4.813467	-1	AREA STUDIES	501	...	
5	AWq8Wt0n9WDJ8JPhaXCC	utexas	4.479859	-1	AREA STUDIES	501	...	
6	AWq8Wt0n9WDJ8JPhaXCH	utexas	4.479859	-1	AREA STUDIES	501	...	
7	AWq8Wt0n9WDJ8JPhaXCD	utexas	4.443958	-1	AREA STUDIES	501	...	
8	AWq8Wt0n9WDJ8JPhaXCF	utexas	4.443958	-1	AREA STUDIES	501	...	



Summarizing and Sorting

Get fancier with your searching!

Sorting in Query

Query, then sort the output

```
{ "query":  
  { "bool" : {  
    "must" : [ { "range": { "p50_earnings": { "gte" : 75000 } } }  
              , { "match": { "institution_id": "3658" } }  
              , { "range" : { "cellcount" : { "gt" : 0 } } } ]  
  }  
},  
"sort": {  
  "ciptitle.raw": "asc"  
}  
}
```

Result Sample

_index	_source.cellcount	_source.ciptitle	_source.deglevel	_source.institution_name	_source.p25_earnings	_source.p50_earnings
0	utexas	450	ACCOUNTING AND RELATED SERVICES	Baccalaureate	... UNIVERSITY OF TEXAS-AUSTIN	75365.58
1	utexas	401	ACCOUNTING AND RELATED SERVICES	Baccalaureate	... UNIVERSITY OF TEXAS-AUSTIN	93491.84
2	utexas	525	ACCOUNTING AND RELATED SERVICES	Baccalaureate	... UNIVERSITY OF TEXAS-AUSTIN	68325.54
3	utexas	536	ACCOUNTING AND RELATED SERVICES	Baccalaureate	... UNIVERSITY OF TEXAS-AUSTIN	70508.76
4	utexas	317	ACCOUNTING AND RELATED SERVICES	Baccalaureate	... UNIVERSITY OF TEXAS-AUSTIN	83546.25
5	utexas	131	AEROSPACE, AERONAUTICAL AND ASTRONAUTICAL ENGI...	Baccalaureate	... UNIVERSITY OF TEXAS-AUSTIN	83546.25
6	utexas	175	AEROSPACE, AERONAUTICAL AND ASTRONAUTICAL ENGI...	Baccalaureate	... UNIVERSITY OF TEXAS-AUSTIN	63792.66
7	utexas	193	AEROSPACE, AERONAUTICAL AND ASTRONAUTICAL ENGI...	Baccalaureate	... UNIVERSITY OF TEXAS-AUSTIN	66558.09
8	utexas	77	AEROSPACE, AERONAUTICAL AND ASTRONAUTICAL ENGI...	Baccalaureate	... UNIVERSITY OF TEXAS-AUSTIN	81609.44
9	utexas	75	AEROSPACE, AERONAUTICAL AND ASTRONAUTICAL ENGI...	Baccalaureate	... UNIVERSITY OF TEXAS-AUSTIN	67923.63

Things to notice:

- ciptitle field is sorted alphabetically
- Only UT-Austin is shown
- p50_earnings are all above \$75,000 as requested

Summarizing in Query

Summarize one field:

```
{  
  "aggs": {  
    "common_majors": {  
      "terms": {  
        "field": "ciptitle.raw"  
      }  
    }  
  }  
}
```

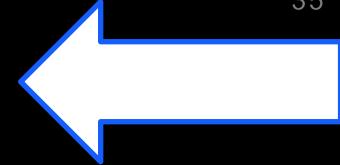
Create a new field called common_majors, which sums up the unique values in ciptitle.raw.

Python Tip! Extract the correct piece of your query result: res['aggregations']['common_majors']['buckets']

Result Sample

Produces:

	doc_count	key
0	153	HISTORY
1	149	HEALTH PROFESSIONS AND RELATED PROGRAMS
2	130	BIOLOGICAL AND BIOMEDICAL SCIENCES
3	93	CLINICAL/MEDICAL LABORATORY SCIENCE/RESEARCH A...
4	90	MATHEMATICS
5	80	LIBERAL ARTS AND SCIENCES, GENERAL STUDIES AND...
6	78	PHYSICAL SCIENCES
7	77	CHEMISTRY
8	75	HEALTH AND PHYSICAL EDUCATION/FITNESS
9	75	POLITICAL SCIENCE AND GOVERNMENT



If you're following along, clean up

- At Terminal: `./supporting_materials/cleanup_local.sh`

This shuts down the docker container, destroying our demo database – but you can create it again just by going back to the beginning.



Further Reading



Explore More about Elasticsearch!

37

ES Query Language

- <http://elasticsearch-cheatsheet.jolicode.com/>
- https://elasticsearch-dsl.readthedocs.io/en/latest/search_dsl.html
- https://www.elastic.co/guide/en/elasticsearch/reference/current/_introducing_the_query_language.html
- <https://www.elastic.co/guide/en/elasticsearch/reference/6.4/query-dsl-bool-query.html>
- <https://www.elastic.co/guide/en/elasticsearch/reference/6.4/query-filter-context.html>
- <https://logz.io/blog/elasticsearch-queries/>

github.com/skirmer/elastic_analytics
www.stephaniekirmer.com
@data_stephanie

Library Docs

- <https://elasticsearch-py.readthedocs.io/en/master/index.html>
- <https://github.com/ropensci/elastic>
- <https://github.com/UptakeOpenSource/uptasticsearch> – Make contributions, the packages are always improving!

Data Credit:

The data being used in this tutorial is from data.world, and comes out of the hard work done by Annie Millerbernd of the San Antonio Express-News. You can learn more about it and see the original dataset here: <https://data.world/amillerbernd/ut-system-post-grad-earnings>

