

Integrating Elasticsearch into Analytics Workflows

REV2

May 23, 2019

STEPHANIE KIRMER



@data_stephanie

github.com/skirmer/elastic_analytics

AGENDA

Introducing Elasticsearch
Libraries for R and Python
Querying and Filtering
Summarizing
Further Reading



Introducing Elasticsearch

Overview


- Part of a family of data storage options called **NoSQL**
 - **NOT** the same as tabular or SQL style data storage
- Optimized for fast and **powerful searching**
- **Scales to “big data”**– but usable for small projects
- **Open source tool**

*Sometimes abbreviated “ES” – don’t get confused by this!



Visualizing Elasticsearch Storage


- **Cluster** = group of nodes
 - **Master Node** = central brain
 - Home of search capabilities
- **Data Node** = where data is stored
 - Where master node looks for documents



Visualizing Elasticsearch Storage

Data Architecture

- Data is divided into **indices**
- **Indices :**
 - Are user-defined groupings of data with some commonality
 - can live on one node, or
 - can be “sharded” and broken across nodes
 - can be duplicated on different nodes





Tabular Data vs Document-Based Data

NoSQL is a different paradigm for thinking about data.

	institution_id	institution_name	deglevl_code	deglevel	degcip_4dig	ciptitle	
1	3599	UNIVERSITY OF TEXAS - RIO GRANDE VALLEY	3	Baccalaureate	301	NATURAL RESOURCES CONSERVATION AND RESEARCH	
2	3599	UNIVERSITY OF TEXAS - RIO GRANDE VALLEY	3	Baccalaureate	301	NATURAL RESOURCES CONSERVATION AND RESEARCH	
3	3599	UNIVERSITY OF TEXAS - RIO GRANDE VALLEY	3	Baccalaureate	501	AREA STUDIES	
4	3599	UNIVERSITY OF TEXAS - RIO GRANDE VALLEY	3	Baccalaureate	501	AREA STUDIES	
5	3599	UNIVERSITY OF TEXAS - RIO GRANDE VALLEY	3	Baccalaureate	501	AREA STUDIES	
6	3599	UNIVERSITY OF TEXAS - RIO GRANDE VALLEY	3	Baccalaureate	501	AREA STUDIES	
7	3599	UNIVERSITY OF TEXAS - RIO GRANDE VALLEY	3	Baccalaureate	501	AREA STUDIES	
8	3599	UNIVERSITY OF TEXAS - RIO GRANDE VALLEY	3	Baccalaureate	501	AREA STUDIES	

```
{ "_index": "utexas",  
  "_type": "data",  
  "_id": "AWbU6WJiWX1fgzrfh4p1",  
  "_score": 1.0,  
  "_source": {  
    "institution_id": 3599,  
    "institution_name": "UNIVERSITY OF TEXAS - RIO GRANDE VALLEY",  
    "deglevl_code": 3,  
    "deglevel": "Baccalaureate",  
    "degcip_4dig": 901,  
    "ciptitle": "COMMUNICATION AND MEDIA STUDIES",  
    "grad_cohort": 2007,  
    "grad_cohort_label": "2007-2009",  
    "year_postgrad": 1,  
    "p25_earnings": 26518.57,  
    "p50_earnings": 42166.31,  
    "p75_earnings": 50439,  
    "system": "utsys",  
    "cellcount": 70 } }
```



Why Use Elasticsearch?



Safe

- Copying your data easily and conveniently (via replicas) = if a node fails, your data is safe



Fast

- ES can search in parallel on multiple nodes and replicas, and find your data faster



Scalable

- Once you establish your ES database, you can add nodes and allow your database to grow



Open Source

- Free to use at small scale, substantial documentation, community support



Follow Along!

When you see this arrow,
you can try it out yourself!

System Requirements:

- Docker installed and running
- Cloned: www.github.com/skirmer/elastic_analytics

Setup Steps (see the README to copy/paste)

- Get into the top level of the cloned repo
- At Terminal:
 1. `./supporting_materials/setup_texas.sh 5.5`
 2. `curl -X POST 'http://localhost:9200/utexas/_bulk' -H 'Content-Type: application/json' --data-binary @supporting_materials/ut_data.json`

Start up R/Rstudio or Python as you prefer, and run further commands from there.



Libraries for R and Python

Choosing the right tool for your
needs



Library Characteristics

Library	Returns	Query Language	Supports Authentication	R	Python
uptasticsearch	Tabular	Required	✗	●	●
elastic	JSON	Supported, not required	●	●	✗
elasticsearch-py	JSON	Supported, not required	●	✗	●

KEY CONSIDERATIONS

Secure Authentication
Do you need to securely log in?


Output Format
Do you mind handling JSON output?

Query Construction
Is writing query language a barrier?



R Options

12



Uptasticsearch (R)

```
test_up <- uptasticsearch::es_search(  
  es_host = "http://localhost:9200"  
  , es_index = "utexas"  
  , query_body = query_string  
  , max_hits = 10  
  , size = 10)
```

```
query_string <- '{"query": {"match_all":{}}}'
```

Elastic (R)

```
elastic::connect(es_host =  
  "http://localhost:9200")  
  
test_e <- elastic::Search(index =  
  "utexas"  
    , body = query_string  
    , size = 10  
    , raw = TRUE)  
  
test_e2 <-  
jsonlite::fromJSON(test_e)$hits$hits
```


Non-Query Language Option:

```
test_e <- elastic::Search(  
  index = "utexas"  
  , q = "grad_cohort:*"   
  , size = 10  
  , raw = TRUE)
```



Python Options

13



Uptasticsearch (Py)

```
import json
import uptasticsearch

es_search(
    es_host="http://localhost:9200",
    query_body=query_string,
    max_hits = 10,
    es_index="utexas"
)
```

```
query_text = {"query": {"match_all": {}}}
query_string = '{"query": {"match_all": {}}}'
```

Elasticsearch-py (Py)

```
from elasticsearch import Elasticsearch

es = Elasticsearch(['http://localhost:9200'])

res = es.search(
    index="utexas",
    body= query_text
)

res['hits']['hits']
```

Non-Query Language Option (Elasticsearch_dsl):

```
from elasticsearch_dsl import Search

res2 = Search(using = es).query("match",
    _index = 'utexas').execute()

res2.to_dict()['hits']['hits']
```



Querying and Filtering

Get what you need out of your
database



Query Language Crash Course

Elastic Query DSL (domain specific language): a JSON-style syntax built to interact with ES databases.


Why use query language?

Consistency across interfaces and media

Precision and power in search, filtering, and aggregating – ES was built to work with this.

Downsides?

It's sometimes hard to work with – idiosyncratic rules of syntax.



Identifying Available Fields

R:

```
uptasticsearch::get_fields(es_host = "http://localhost:9200",  
es_indices = "utexas")
```

At Command Line:

```
curl http://localhost:9200/utexas/_mapping > fields.json
```



Constructing a Basic Query

Return all records :

```
{  
  "query": { "match_all": { } }  
}
```

```
query_text = {"query": {"match_all": {}}}  
query_string = '{"query": {"match_all": {}}}'
```

Remember this from earlier!
All the queries we look at can be
passed to R or Python this way.





Constructing a Basic Query

Return all records :

```
{  
  "query": { "match_all": { } }  
}
```



Constructing a Basic Query

Match one field :

```
{  
  "query": { "match": { "ciptitle.raw": "AREA STUDIES" } }  
}
```

Constructing a Basic Query

Match one field AND Greater Than one field :

```
{  
  "query":  
    { "bool" : {  
      "must" : { "match": { "ciptitle.raw": "AREA STUDIES" } }  
      , "must" : { "range" : { "cellcount" : { "gte" : 0 } } }  
    } }  
}
```



Constructing a Basic Query

Match two fields AND Greater Than one field :

```
{
  "query":
    { "bool" : {
      "must" : [ { "match": { "ciptitle.raw": "AREA STUDIES" } }
        , { "match": { "institution_id": "3599" } }
        , { "range" : { "cellcount" : { "gte" : 0 } } } ]
    } }
}
```



Some Other Querying Options

22

match_phrase

Match a set of words all together.

exists

Supply a field, returns documents that have at least one non-null value in the original field.

wildcard

Pass a string with a wildcard anywhere – but be careful, it can be a slow search!

filter

Just like “must” except without scoring – we’ll talk about this in a moment.

must_not

Instead of “must” – use to omit records with a word or phrase.

This is just a small sample- ES query language offers many very powerful search options!


Query vs Filter

Scoring Results

ES queries can provide a **numeric score** indicating how well the document meets the criteria given.

When you use “**query**” at the beginning of the query, you get a score returned alongside your results.

When you use “**filter**”, Elasticsearch does not score the results on the given criteria.



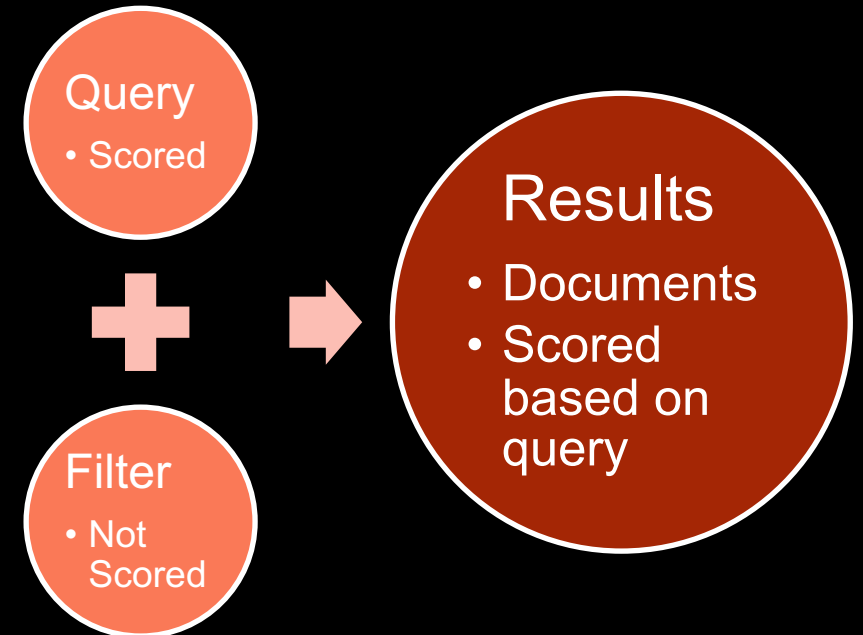
Query vs Filter

Example

Do we want scores returned for this search? Yes.

```
{ "query":
  { "bool": {
    "must": [
      { "match": { "ciptitle.raw": "AREA STUDIES" } }
      , { "match": { "deglevel": "Baccalaureate" } }
    ] ,
    "filter": [
      { "match": { "institution_id": "3599" } }
      , { "range": { "cellcount" : { "gte" : 0 } } }
    ]
  } }
}
```

Do we want the scores to include these criteria? NO.



Query vs Filter

Example

Same query, first with one criterion scored (2 in filter) and second with all 3 criteria scored.

	_index	_type	_id	_score	institution_id	institution_name
1	utexas	data	AWbk3KyPrRbatlpY56KW	5.159055	3599	UNIVERSITY OF TEXAS – RIO GRA
2	utexas	data	AWbk3KyPrRbatlpY56KU	4.790120	3599	UNIVERSITY OF TEXAS – RIO GRA
3	utexas	data	AWbk3KyPrRbatlpY56KS	4.699614	3599	UNIVERSITY OF TEXAS – RIO GRA
4	utexas	data	AWbk3KyPrRbatlpY56KX	4.699614	3599	UNIVERSITY OF TEXAS – RIO GRA

	_index	_type	_id	_score	institution_id	institution_name
1	utexas	data	AWbk3KyPrRbatlpY56KW	7.159055	3599	UNIVERSITY OF TEXAS – RIO GRAND
2	utexas	data	AWbk3KyPrRbatlpY56KU	6.790120	3599	UNIVERSITY OF TEXAS – RIO GRAND
3	utexas	data	AWbk3KyPrRbatlpY56KS	6.699614	3599	UNIVERSITY OF TEXAS – RIO GRAND
4	utexas	data	AWbk3KyPrRbatlpY56KX	6.699614	3599	UNIVERSITY OF TEXAS – RIO GRAND



Summarizing

Get fancier with your searching!

Summarizing in Query

Match one field, Summarize one field:

```
{  
  "query": {"match": {"institution_id": "3599" }} ,  
  "aggs" : {  
    "common_majors" : {  
      "terms" : {  
        "field" : "ciptitle.raw"  
      }  
    }  
  }  
}
```


Create a new field called `common_majors`, which sums up the unique values in `ciptitle.raw` for this specific search.



Summarizing in Query

Produces:

	common_majors	doc_count
1	HISTORY	12
2	BUSINESS, MANAGEMENT, MARKETING, AND RELATED SUPPORT SERVICES	11
3	EDUCATION	11
4	ANTHROPOLOGY	10
5	CLINICAL/MEDICAL LABORATORY SCIENCE/RESEARCH AND ALLIED PROFESSIONS	10
Showing 1 to 5 of 10 entries		Previous 1 2 Next



If you're following along, clean up

- At Terminal: `./supporting_materials/cleanup_local.sh`

This shuts down the docker container, destroying our demo database – but you can create it again just by going back to the beginning.



Further Reading



ES Query Language

- <http://elasticsearch-cheatsheet.jolicode.com/>
- https://elasticsearch-dsl.readthedocs.io/en/latest/search_dsl.html
- https://www.elastic.co/guide/en/elasticsearch/reference/current/_introducing_the_query_language.html
- <https://www.elastic.co/guide/en/elasticsearch/reference/6.4/query-dsl-bool-query.html>
- <https://www.elastic.co/guide/en/elasticsearch/reference/6.4/query-filter-context.html>

github.com/skirmer/elastic_analytics
www.stephaniekirmer.com
[@data_stephanie](#)

Library Docs

- <https://elasticsearch-py.readthedocs.io/en/master/index.html>
- <https://github.com/ropensci/elastic>
- <https://github.com/UptakeOpenSource/uptasticsearch> – Make contributions, the packages are always improving!

Data Credit:

The data being used in this tutorial is from data.world, and comes out of the hard work done by Annie Millerbernd of the San Antonio Express-News. You can learn more about it and see the original dataset here: <https://data.world/amillerbernd/ut-system-post-grad-earnings>

